

Q1 Machine Learning Project Report: Credit Score Classification

By: Vatsal Sivaratri

9/18/2024

Table of Contents

Part 1 – Statement/Project Goal	3
Part 2 – Description of Dataset	3
Part 3 – Methods and Materials	3
Part 4 – Results	14

1 Statement/Project Goal

The goal of this project is to develop a predictive, robust model(s) for classifying individuals' credit scores. The motivation behind this project is the growing need for transparency and accuracy in evaluating trustworthiness of Credit Card users. A strong relationship between banks and their customers is necessary for a productive economy. The hope is to use past data on credit card holders to predict the likelihood of a good credit score in new customers, by considering a variety of attributes.

2 Description of Dataset

The data was found on [Kaggle](#) from a competition hosted by the website themselves. This dataset has conveniently been split into training and testing datasets. The training dataset contains 100,000 instances and 27 attributes (Dimension – 28). The attributes include both numerical data (e.g., monthly in-hand salary, outstanding debt) and categorical data (e.g., occupation, credit mix). Here's a list of the key attributes:

Age: Age of individual

Occupation: Profession of Individual

Annual_Income: Annual earnings of individual

Monthly_Inhand_Salary: Amount individual receives per month.

Num_Bank_Accounts: Number of bank accounts the individual holds

Outstanding_Debt: Amount of unpaid debt

Credit_Utilization_Ratio: The ratio of the credit limit utilized

Credit_History_Age: The duration of the individual's credit history

Credit_Mix: Information on the mix of credit types (e.g., Good, Bad).

Monthly_Balance: The remaining balance at the end of each month after all payments.

There are some interesting observations, including potential outliers in the age column (like negative values for age), repeated values in columns that are supposed to be unique (SSN), and missing values in some columns.

3 Methods and Materials

Data Preprocessing:

When we first start with our dataset, we have a total of 27 features, 1 class variable, with 100,000 instances.

```

0   ID                100000 non-null object
1   Customer_ID       100000 non-null object
2   Month             100000 non-null object
3   Name              90015  non-null object
4   Age               100000 non-null object
5   SSN               100000 non-null object
6   Occupation        100000 non-null object
7   Annual_Income     100000 non-null object
8   Monthly_Inhand_Salary  84998 non-null float64
9   Num_Bank_Accounts 100000 non-null int64
10  Num_Credit_Card    100000 non-null int64
11  Interest_Rate     100000 non-null int64
12  Num_of_Loan        100000 non-null object
13  Type_of_Loan       88592 non-null object
14  Delay_from_due_date 100000 non-null int64
15  Num_of_Delayed_Payment 92998 non-null object
16  Changed_Credit_Limit 100000 non-null object
17  Num_Credit_Inquiries 98035 non-null float64
18  Credit_Mix         100000 non-null object
19  Outstanding_Debt   100000 non-null object
20  Credit_Utilization_Ratio 100000 non-null float64
21  Credit_History_Age 90970  non-null object
22  Payment_of_Min_Amount 100000 non-null object
23  Total_EMI_per_month 100000 non-null float64
24  Amount_invested_monthly 95521 non-null object
25  Payment_Behaviour  100000 non-null object
26  Monthly_Balance    98800 non-null object
27  Credit_Score       100000 non-null object
dtypes: float64(4), int64(4), object(20)
memory usage: 21.4+ MB

```

1. **Removing any Overtly Useless Columns:** Ex. A person's SSN is useless in predicting their credit score without any other relevant information given by it (and even then it shouldn't be needed as a feature itself)=> Go through all features and see which ones aren't productive and need to be removed.

```

useless_columns = ['ID', 'Customer_ID', 'Name', 'SSN']
data.drop(useless_columns, axis=1, inplace=True)

```

2. **Handling Underscore Issue:** In this dataset, there appears to be a continued problem of instances have an extra underscore, which prevents further analysis.

```

data['Annual_Income'].unique()
[308] ✓ 0.0s Python
... array(['19114.12', '34847.84', '34847.84_', ..., '20002.88', '39628.99',
        '39628.99_'], dtype=object)

```

```
def remove_trailing_underscore(df, column_name):
    def clean_value(value):
        if isinstance(value, str) and value.endswith('_'):
            return value[:-1]
        return value

    df[column_name] = df[column_name].apply(clean_value)

    return df

for column in data.columns:
    data = remove_trailing_underscore(data, column)
```

- 3. Outlier Detection and Removal:** Certain variables like “Age” and “Monthly_Balance” can have extreme values that are outliers and bad representations of other cars with similar specifications. These will need to be identified and removed, both manually and with IQR.

Graph distributions of features to see any apparent outliers:

```
import matplotlib.pyplot as plt
import seaborn as sns

def plot_feature_distributions(data, plot_type='hist'):
    num_columns = len(data.columns)

    plt.figure(figsize=(16, num_columns * 4))

    for i, column in enumerate(data.columns, 1):
        plt.subplot(num_columns, 1, i)

        if plot_type == 'hist':
            sns.histplot(data[column], kde=True, bins=100)
            plt.title(f'Distribution of {column}')
        elif plot_type == 'box':
            sns.boxplot(x=data[column])
            plt.title(f'Boxplot of {column}')

    plt.tight_layout()

    plt.show()

plot_feature_distributions(data, plot_type='hist')
```

Output:

https://drive.google.com/file/d/1r5UP0r1FlC_KLzjFjZdry6x4WfDCn0qW/view?usp=sharing

Remove Filler Outliers: Datasets like these often have people inputting obviously wrong data as placeholders. These methods take care of those issues.

```
def process_age(data):
    age = pd.to_numeric(data['Age'], errors='coerce')

    wrong_ages = age[(age == -500) | (age > 995) | (age < 18)].index

    data = data.drop(wrong_ages)

    data['Age'] = pd.to_numeric(data['Age'], errors='coerce')

    return data
data = process_age(data)

def process_ir(data):
    ir = pd.to_numeric(data['Interest_Rate'], errors='coerce')

    wrong_irs = ir[(ir < 0) | (ir > 100)].index

    data = data.drop(wrong_irs)

    data['Interest_Rate'] = pd.to_numeric(data['Interest_Rate'],
errors='coerce')

    return data
data = process_ir(data)

def convert_credit_history_age_to_months(data):
    def convert_to_months(value):
        if pd.isna(value):
            return None
        parts = value.split('and')
        years = int(parts[0].strip().split()[0])
        months = int(parts[1].strip().split()[0])
```

```

        return years * 12 + months

    data['Credit_History_Age'] =
data['Credit_History_Age'].apply(convert_to_months)

    data['Credit_History_Age'].fillna(data['Credit_History_Age'].mean(),
inplace=True)

    return data
data = convert_credit_history_age_to_months(data)

def process_delayed_payment(data):
    delayed_payment = pd.to_numeric(data['Num_of_Delayed_Payment'],
errors='coerce')

    wrong_irs = delayed_payment[((delayed_payment < 0))].index

    data = data.drop(wrong_irs)

    data['Num_of_Delayed_Payment'] =
pd.to_numeric(data['Num_of_Delayed_Payment'], errors='coerce')

    return data
data = process_delayed_payment(data)

def process_amount_invested_monthly(data):

    data['Amount_invested_monthly'] =
data['Amount_invested_monthly'].replace('__10000__', pd.NA)

    data['Amount_invested_monthly'] =
pd.to_numeric(data['Amount_invested_monthly'], errors='coerce')

    return data

data = process_amount_invested_monthly(data)

def process_payment_behavior(data):

```

```

data['Payment_Behaviour'] = data['Payment_Behaviour'].replace('@9#%8', '0')

return data
data = process_payment_behavior(data)

```

For quantitative continuous features, also remove outliers with IQR:

```

features_to_process = ['Age', 'Annual_Income', 'Monthly_Inhand_Salary',
                        'Outstanding_Debt', 'Credit_Utilization_Ratio',
                        'Total_EMI_per_month', 'Amount_invested_monthly']

for feature in features_to_process:
    data[feature] = pd.to_numeric(data[feature], errors='coerce')

    data[feature].fillna(data[feature].mean(), inplace=True)

    Q1 = data[feature].quantile(0.25)
    Q3 = data[feature].quantile(0.75)
    IQR = Q3 - Q1

    data = data[~((data[feature] < (Q1 - 1.5 * IQR)) | (data[feature] > (Q3 +
1.5 * IQR)))]

```

4. **Fix Loan Feature Representations:** Loans are a very important feature in determining what kind of credit score a person has. However, the way loans are currently represented in the dataset combines all of them as one feature, making it hard to represent meaningfully to the model (Too many combinations of loan types). Thus, split this feature into 10 other features, 1 for each possible category:

```

def add_loan_type_columns(data):
    unique_loan_types = ['Auto Loan', 'Credit-Builder Loan', 'Debt Consolidation
Loan', 'Home Equity Loan', 'Mortgage Loan',
                        'No Loan', 'Not Specified', 'Payday Loan', 'Personal Loan',
                        'Student Loan']

    for loan_type in unique_loan_types:

        cleaned_loan_type = loan_type.replace(' ', '_').replace('-',
'_').lower()

        data[cleaned_loan_type] = data['Type_of_Loan'].apply(lambda x:
x.count(loan_type) if isinstance(x, str) else 0)
        data.drop(['Type_of_Loan'], axis=1, inplace=True)

```



```

return data
data = add_loan_type_columns(data)

```

5. **Handle Missing Data:** As mentioned before, several columns contain missing values. How to fill them depends on the feature being dealt with. For example, the column “Monthly_Inhand_Salary” has missing values, and since this is a numerical, continuous feature, we may best fill them using the mean of the column. For “Num_Credit_Inquiries”, “Occupation” and other discrete quantitative or qualitative features, values will be filled with mode:

```

mean_columns = ['Monthly_Inhand_Salary', 'Amount_invested_monthly',
'Monthly_Balance', 'Changed_Credit_Limit']
mode_columns = ['Num_of_Delayed_Payment', 'Num_Credit_Inquiries']
categorical_mode = ['Payment_Behaviour']

for column in mean_columns:
    data[column] = pd.to_numeric(data[column], errors='coerce')

    data[column].fillna(data[column].mean(), inplace=True)
for column in mode_columns:
    data[column] = pd.to_numeric(data[column], errors='coerce')

    data[column].fillna(data[column].mode()[0], inplace=True)

```

6. **Encoding Categorical Variables:** Columns like “Occupation” and “Credit_Mix” need to be encoded as numerical values for the model. This might involve one-hot encoding or label encoding.

```

toEncode = ['Month', 'Occupation', 'Credit_Mix', 'Payment_of_Min_Amount',
'Payment_Behaviour', 'Credit_Score']
label_encoder = LabelEncoder()
for column in toEncode:
    data[column] = label_encoder.fit_transform(data[column])

```

After preprocessing the data, we went from 100,000 instances and 27 features to 74652 instances and 32 features.

Feature Selection: 4 Feature Selection Algorithms will be selected and applied onto the data through WEKA, and one data subset will be selected separately.

Method 1: CorrelationAttributeEval

```

0.18151 12 Changed_Credit_Limit
0.17729 11 Num_of_Delayed_Payment
0.17391 18 Payment_of_Min_Amount
0.16813 6 Num_Bank_Accounts
0.13492 14 Credit_Mix

```

0.11204 8 Interest_Rate
 0.10648 7 Num_Credit_Card
 0.08069 10 Delay_from_due_date
 0.05262 13 Num_Credit_Inquiries
 0.04484 9 Num_of_Loan
 0.02833 31 personal_loan
 0.0261 15 Outstanding_Debt
 0.02551 25 debt_consolidation_loan
 0.02287 21 Payment_Behaviour
 0.02244 23 auto_loan
 0.02057 26 home_equity_loan
 0.02003 24 credit_builder_loan
 0.01703 27 mortgage_loan
 0.01683 30 payday_loan
 0.01486 29 not_specified
 0.00501 3 Occupation
 0.00482 32 student_loan
 0.00318 1 Month
 0 28 no_loan
 -0.00864 16 Credit_Utilization_Ratio
 -0.01213 19 Total_EMI_per_month
 -0.02063 20 Amount_invested_monthly
 -0.03455 5 Monthly_Inhand_Salary
 -0.03914 4 Annual_Income
 -0.04084 22 Monthly_Balance
 -0.04399 2 Age
 -0.08273 17 Credit_History_Age

Cut-off: 0.10 => Selected Features:

0.18151 12 Changed_Credit_Limit
 0.17729 11 Num_of_Delayed_Payment
 0.17391 18 Payment_of_Min_Amount
 0.16813 6 Num_Bank_Accounts
 0.13492 14 Credit_Mix
 0.11204 8 Interest_Rate
 0.10648 7 Num_Credit_Card

Method 2: ReliefAttributeEval (Ranker):

0.024477 2 Age
 0.017906 7 Num_Credit_Card
 0.017332 6 Num_Bank_Accounts
 0.017219 8 Interest_Rate
 0.015441 15 Outstanding_Debt
 0.014112 19 Total_EMI_per_month
 0.014064 3 Occupation
 0.0127 17 Credit_History_Age

0.011524 11 Num_of_Delayed_Payment
 0.01107 4 Annual_Income
 0.010864 10 Delay_from_due_date
 0.009616 13 Num_Credit_Inquiries
 0.00737 12 Changed_Credit_Limit
 0.006013 9 Num_of_Loan
 0.005844 5 Monthly_Inhand_Salary
 0.005668 32 student_loan
 0.005657 30 payday_loan
 0.005379 29 not_specified
 0.005245 27 mortgage_loan
 0.005222 24 credit_builder_loan
 0.00521 23 auto_loan
 0.005032 31 personal_loan
 0.004947 25 debt_consolidation_loan
 0.0048 26 home_equity_loan
 0 28 no_loan
 -0.000387 22 Monthly_Balance
 -0.009254 20 Amount_invested_monthly
 -0.018253 16 Credit_Utilization_Ratio
 -0.021545 18 Payment_of_Min_Amount
 -0.023212 14 Credit_Mix
 -0.032999 21 Payment_Behaviour
 -0.060802 1 Month

Cut-off: 0.02 => Selected Features:

0.024477 2 Age
 -0.021545 18 Payment_of_Min_Amount
 -0.023212 14 Credit_Mix
 -0.032999 21 Payment_Behaviour
 -0.060802 1 Month

Method 3: ClassifierAttributeEval (Ranker)

0.1879941 4 Annual_Income
 0.1812442 15 Outstanding_Debt
 0.1583674 19 Total_EMI_per_month
 0.1448428 5 Monthly_Inhand_Salary
 0.0667689 14 Credit_Mix
 0.0327237 8 Interest_Rate
 0.0273498 7 Num_Credit_Card
 0.0242182 18 Payment_of_Min_Amount
 0.0229147 11 Num_of_Delayed_Payment
 0.0222949 6 Num_Bank_Accounts

0.0188205 10 Delay_from_due_date
 0.0187684 12 Changed_Credit_Limit
 0.0048473 13 Num_Credit_Inquiries
 0.0016956 2 Age
 0.001023 9 Num_of_Loan
 0.0004473 1 Month
 0.0003747 31 personal_loan
 0.000297 25 debt_consolidation_loan
 0.0002325 23 auto_loan
 0.0002311 17 Credit_History_Age
 0.000171 21 Payment_Behaviour
 0.0001605 29 not_specified
 0.0001526 3 Occupation
 0.0001436 27 mortgage_loan
 0.000138 26 home_equity_loan
 0.0001307 24 credit_builder_loan
 0.0000864 30 payday_loan
 0.0000121 32 student_loan
 0 28 no_loan
 -0.1479401 20 Amount_invested_monthly
 -0.1575615 22 Monthly_Balance
 -0.1602462 16 Credit_Utilization_Ratio

Cutoff: 0.10 => Selected Features:

0.1879941 4 Annual_Income
 0.1812442 15 Outstanding_Debt
 0.1583674 19 Total_EMI_per_month
 0.1448428 5 Monthly_Inhand_Salary
 -0.1479401 20 Amount_invested_monthly
 -0.1575615 22 Monthly_Balance
 -0.1602462 16 Credit_Utilization_Ratio

Method 4: CfsSubsetEval (GreedyStepwise)

Algorithm's Selected Features:

Num_Bank_Accounts
 Num_of_Delayed_Payment
 Changed_Credit_Limit
 Credit_Mix
 Payment_of_Min_Amount

Method 5 (Selecting Subset Manually):

'Annual_Income','Monthly_Inhand_Salary','Num_of_Loan',
 'Outstanding_Debt','Credit_Utilization_Ratio','Credit_History_Age',

'Num_Credit_Card','Num_of_Delayed_Payment','Interest_Rate','Total_EMI_per_month'

I selected these features from research on the biggest influencers of Credit Score combined with my own intuition.

Models: Using Decision Tree, Random Forest, SVM, Logistic Regression, and Naive Bayes.

For Metrics: Accuracy to evaluate overall performance, ROC-AUCs to see how well the model tells the difference between classes, and TP and FP rates.

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

models = {
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Naive Bayes": GaussianNB(),
    "SVM": SVC(probability=True)
}

results = {}

for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_pred_proba = model.predict_proba(X_test) if hasattr(model, 'predict_proba') else
None

    accuracy = accuracy_score(y_test, y_pred)

    cm = confusion_matrix(y_test, y_pred)

    tpr = {}
    fpr = {}
    for i in range(cm.shape[0]):
        tp = cm[i, i]
        fn = cm[i, :].sum() - tp
        fp = cm[:, i].sum() - tp
        tn = cm.sum() - (tp + fn + fp)

        tpr[i] = tp / (tp + fn) if (tp + fn) != 0 else 0
```

```

        fpr[i] = fp / (fp + tn) if (fp + tn) != 0 else 0

    roc_auc = roc_auc_score(y_test, y_pred_proba, multi_class='ovr') if y_pred_proba is
not None else 'N/A'

    results[model_name] = {
        'Accuracy': accuracy,
        'Confusion Matrix': cm,
        'True Positive Rate': tpr,
        'False Positive Rate': fpr,
        'ROC-AUC': roc_auc
    }

for model_name, metrics in results.items():
    print(f"Model: {model_name}")
    print(f"Accuracy: {metrics['Accuracy']}")
    print(f"Confusion Matrix: \n{metrics['Confusion Matrix']}")
    print(f"True Positive Rate (TPR): {metrics['True Positive Rate']}")
    print(f"False Positive Rate (FPR): {metrics['False Positive Rate']}")
    print(f"ROC-AUC: {metrics['ROC-AUC']}\n")

best_model = max(results, key=lambda x: results[x]['Accuracy'])
print(f"The best model is: {best_model}")

```

4 Results

CorrelationAttributeEval

Model	Accuracy	ROC-AUC	TPR_Class_0	FPR_Class_0	TPR_Class_1	FPR_Class_1	TPR_Class_2	FPR_Class_2
Decision Tree	0.688	0.755	0.624	0.09	0.682	0.149	0.71	0.285
Random Forest	0.738	0.871	0.644	0.07	0.699	0.103	0.787	0.284
Logistic Regression	0.637	0.778	0.365	0.06	0.474	0.096	0.806	0.533
Naive Bayes	0.629	0.769	0.754	0.186	0.612	0.151	0.6	0.236

SVM	0.659	0.793	0.358	0.057	0.498	0.073	0.836	0.527
-----	-------	-------	-------	-------	-------	-------	-------	-------

ReliefAttributeEval

Model	Accuracy	ROC-AUC	TPR_Class_0	FPR_Class_0	TPR_Class_1	FPR_Class_1	TPR_Class_2	FPR_Class_2
Decision Tree	0.539	0.677	0.54	0.148	0.438	0.216	0.593	0.402
Random Forest	0.554	0.709	0.478	0.118	0.39	0.181	0.664	0.477
Logistic Regression	0.528	0.674	0.068	0.046	0.232	0.11	0.825	0.778
Naive Bayes	0.556	0.714	0.135	0.028	0.273	0.112	0.834	0.747
SVM	0.544	0.718	0	0	0	0	1	1

ClassifierAttributeEval

Model	Accuracy	ROC-AUC	TPR_Class_0	FPR_Class_0	TPR_Class_1	FPR_Class_1	TPR_Class_2	FPR_Class_2
Decision Tree	0.697	0.751	0.615	0.083	0.687	0.126	0.726	0.313
Random Forest	0.719	0.861	0.378	0.033	0.721	0.104	0.821	0.391
Logistic Regression	0.551	0.668	0.004	0.001	0.284	0.111	0.859	0.807
Naive Bayes	0.544	0.66	0.233	0.133	0.588	0.18	0.614	0.474
SVM	0.544	0.69	0.006	0.002	0.231	0.098	0.874	0.841

CfsSubsetEval

Model	Accuracy	ROC-AUC	TPR_Class_0	FPR_Class_0	TPR_Class_1	FPR_Class_1	TPR_Class_2	FPR_Class_2
Decision Tree	0.638	0.716	0.57	0.101	0.63	0.184	0.662	0.32
Random Forest	0.635	0.794	0.521	0.089	0.577	0.17	0.701	0.37
Logistic Regression	0.588	0.724	0.276	0.066	0.36	0.099	0.805	0.625
Naive Bayes	0.59	0.731	0.662	0.164	0.474	0.151	0.63	0.361
SVM	0.603	0.741	0.386	0.094	0.304	0.056	0.829	0.608

My Feature Selection:

Model	Accuracy	ROC-AUC	TPR_Class_0	FPR_Class_0	TPR_Class_1	FPR_Class_1	TPR_Class_2	FPR_Class_2
Decision Tree	0.727	0.774	0.644	0.074	0.714	0.111	0.759	0.287
Random Forest	0.793	0.913	0.706	0.045	0.801	0.091	0.815	0.228
Logistic Regression	0.569	0.731	0.16	0.04	0.353	0.116	0.809	0.689
Naive Bayes	0.587	0.751	0.757	0.256	0.69	0.167	0.481	0.173
SVM	0.543	0.69	0.006	0.002	0.23	0.099	0.874	0.842



