

Credit Score Classification Using Machine Learning

By: Vatsal Sivaratri

Problem Statement and Motivation

Goal: Develop a predictive, robust model(s) for classifying individuals' credit scores.

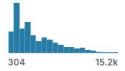
Motivation: There is a growing need for transparency and accuracy in evaluating trustworthiness of Credit Card users. A strong relationship between banks and their customers is necessary for a productive economy. The hope is to use past data on credit card holders to predict the likelihood of a good credit score in new customers, by considering a variety of attributes.

Dataset Source & Description

- **Source:** [Kaggle Card ScoreDataset](#)
- **Key Details:** Training dataset consists of 100,000 instances and 27 features, including both numerical and categorical.
- **Features:** Age, Occupation, Annual_Income, Monthly_Inhand_Salary, Outstanding_Debt, Credit_Mix, etc.
- Data was already split into training and testing set, but additional preprocessing was required to clean and handle inconsistencies. Only training set was used as it had enough samples.
- 3 different classes of Credit_Score: Good, Poor, and Standard

Data Types and Important Features

- **Numerical Data:** Features like Annual_Income, Outstanding_Debt, and Credit_Utilization_Ratio
- **Categorical Data:** Features like Occupation, Credit_Mix, and Payment_Behaviour.
- Many of the features were annoyingly represented as objects, meaning they had to be converted later in Python for further analysis
- Understanding types of features were crucial in determining the preprocessing techniques.

Detail Compact Column						28 of 28 columns
About this file						Add Suggestion
This is the training file						
Age	SSN	Occupation	Annual_Income	Monthly_Inhand...	Nur	
Represents the age of the person	Represents the social security number of a person	Represents the occupation of the person	Represents the annual income of the person	Represents the monthly base salary of a person	Repres of ban person	
38 28 Other (94338)	3% 3% 94% #F%\$D@*48 004-07-5839 Other (94420)	6% 0% 94% _____ Lawyer Other (86363)	7% 7% 86% 18940 unique values		-1	
23	821-00-0265	Scientist	19114.12	1824.8433333333328	3	
23	821-00-0265	Scientist	19114.12		3	
~500	821-00-0265	Scientist	19114.12		3	
23	821-00-0265	Scientist	19114.12		3	
23	821-00-0265	Scientist	19114.12	1824.8433333333328	3	

Initial Issues Found in Data

- **Outliers:** values such as age had anomalies like negative or extremely large values (i.e. -500 and >8900 for age)
 - These are a part of any legitimate dataset, and represent a portion of people who chose not to answer truthfully or were incorrectly recorded
- **Missing Data:** Several columns such as Monthly_Inhand_Salary had missing values that could skew the predictions if left unaddressed
- Trailing underscores were found in several columns

```
0 ID 100000 non-null object
1 Customer_ID 100000 non-null object
2 Month 100000 non-null object
3 Name 90015 non-null object
4 Age 100000 non-null object
5 SSN 100000 non-null object
6 Occupation 100000 non-null object
7 Annual_Income 100000 non-null object
8 Monthly_Inhand_Salary 84998 non-null float64
9 Num_Bank_Accounts 100000 non-null int64
10 Num_Credit_Card 100000 non-null int64
11 Interest_Rate 100000 non-null int64
12 Num_of_Loan 100000 non-null object
13 Type_of_Loan 88592 non-null object
14 Delay_from_due_date 100000 non-null int64
15 Num_of_Delayed_Payment 92998 non-null object
16 Changed_Credit_Limit 100000 non-null object
17 Num_Credit_Inquiries 98035 non-null float64
18 Credit_Mix 100000 non-null object
19 Outstanding_Debt 100000 non-null object
20 Credit_Utilization_Ratio 100000 non-null float64
21 Credit_History_Age 90970 non-null object
22 Payment_of_Min_Amount 100000 non-null object
23 Total_EMI_per_month 100000 non-null float64
24 Amount_invested_monthly 95521 non-null object
25 Payment_Behaviour 100000 non-null object
26 Monthly_Balance 98800 non-null object
27 Credit_Score 100000 non-null object
dtypes: float64(4), int64(4), object(20)
memory usage: 21.4+ MB
```

```
data['Annual_Income'].unique()
[308] ✓ 0.0s Python
... array(['19114.12', '34847.84', '34847.84_', ..., '20002.88', '39628.99'
        '39628.99_'], dtype=object)
```

Outlier Detection and Handling

- Each feature was analyzed manually for extreme or blatantly incorrect values with simple python functions like `min()` and `max()`. They were filtered out the Pandas Dataframe with simple logical statements
- For quantitative continuous variables, additional outliers were detected and removed using Interquartile Range (IQR)

```
...
    age = pd.to_numeric(data['Age'], errors='coerce')

    wrong_ages = age[(age == -500) | (age>995) |
                     (age<18)].index

    data = data.drop(wrong_ages)

    data['Age'] = pd.to_numeric(data['Age'],
                               errors='coerce')
...

Q1 = data[feature].quantile(0.25) Q3 =
data[feature].quantile(0.75) IQR = Q3 - Q1 data =
data[~(((data[feature] < (Q1 - 1.5 * IQR)) |
(data[feature] > (Q3 + 1.5 * IQR))))]
```

Missing Data Handling

For numerical columns, missing data was filled using the **mean**. For categorical columns, missing values were filled using the mode.

```
mean_columns = ['Monthly_Inhand_Salary', 'Amount_invested_monthly', 'Monthly_Balance', 'Changed_Credit_Limit']
mode_columns = ['Num_of_Delayed_Payment', 'Num_Credit_Inquiries']
categorical_mode = ['Payment_Behaviour']
```

```
for column in mean_columns:
    data[column] = pd.to_numeric(data[column], errors='coerce')
```

```
    data[column].fillna(data[column].mean(), inplace=True)
```

```
for column in mode_columns:
    data[column] = pd.to_numeric(data[column], errors='coerce')
```

```
    data[column].fillna(data[column].mode()[0], inplace=True)
```

Encoding Categorical Variables

Label Encoding: Categorical variables were converted into numerical format for models to use using Label Encoding.

```
toEncode = ['Month', 'Occupation', 'Credit_Mix', 'Payment_Behaviour', 'Credit_Score']  
label_encoder = LabelEncoder()
```

```
for column in toEncode:  
    data[column] = label_encoder.fit_transform(data[column])
```

***The classes were also converted like this. Now, 'Good' is 0, 'Poor' is 1, and 'Standard' is 2.**

Feature Engineering

Loan Columns: The original dataset had a column for Type_of_Loan that combined multiple loan types (Auto Loan, Credit-Builder Loan, Debt Consolidation Loan, Home Equity Loan, Mortgage Loan, No Loan, Not Specified, Payday Loan, Personal Loan, Student Loan). This was split into multiple binary columns

```
unique_loan_types = ['Auto Loan', 'Mortgage Loan', 'Personal Loan', ...]
for loan_type in unique_loan_types:
    cleaned_loan_type = loan_type.replace(' ', '_').lower()
    data[cleaned_loan_type] = data['Type_of_Loan'].apply(lambda x: loan_type in x if
isinstance(x, str) else False)
```

Then 'Auto Loan, Mortgage Loan' => {...,Auto Loan: 1, Mortgage Loan: 1....}

Feature Selection Techniques

- **CorrelationAttributeEval:** Method evaluates correlation between individual attributes and the class label (Pearson's Correlation Coefficient). Attributes with high correlation to the class are selected (cutoff=0.10)
- **ReliefAttributeEval:** Algorithm estimate feature importance based on how well feature can distinguish between instances that are near each other (cutoff=0.02)
- **ClassifierAttributeEval:** Ranks features based on how well they contribute to classifier performance (I used RandomForest). Features like Annual_Income and Outstanding_Debt got prioritized here.
- **CfsSubsetEval:** Selects group of features that work well together. Identified attributes like Num_Bank_Accounts and Credit_Mix being most relevant for predicting credit scores.
- **Individually Selected Features:**
'Annual_Income','Monthly_Inhand_Salary','Num_of_Loan',
'Outstanding_Debt','Credit_Utilization_Ratio','Credit_History_Age',
'Num_Credit_Card','Num_of_Delayed_Payment','Interest_Rate','Total_EMI_per_month'

CorrelationAttributeEval

Selected Features:

0.18151	12	Changed_Credit_Limit
0.17729	11	Num_of_Delayed_Payment
0.17391	18	Payment_of_Min_Amount
0.16813	6	Num_Bank_Accounts
0.13492	14	Credit_Mix
0.11204	8	Interest_Rate
0.10648	7	Num_Credit_Card

ReliefAttributeEval

Selected Features:

0.024477	2 Age
-0.021545	18 Payment_of_Min_Amount
-0.023212	14 Credit_Mix
-0.032999	21 Payment_Behaviour
-0.060802	1 Month

ClassifierAttributeEval

Selected Features

0.1879941 4 Annual_Income

0.1812442 15 Outstanding_Debt

0.1583674 19 Total_EMI_per_month

0.1448428 5 Monthly_Inhand_Salary

-0.1479401 20 Amount_invested_monthly

-0.1575615 22 Monthly_Balance

-0.1602462 16 Credit_Utilization_Ratio

CfsSubsetEval

Selected Features:

Num_Bank_Accounts

Num_of_Delayed_Payment

Changed_Credit_Limit

Credit_Mix

Payment_of_Min_Amount

Models

Models Evaluated: Decision Tree, Random Forest, Logistic Regression, Naive Bayes, Support Vector Machines (SVM). Evaluated using Accuracy, True Positive Rates (TPR), False Positive Rates (FPR), and ROC-AUCs

```
models = {  
    "Decision Tree": DecisionTreeClassifier(),  
    "Random Forest": RandomForestClassifier(),  
    "Logistic Regression": LogisticRegression(max_iter=1000),  
    "Naive Bayes": GaussianNB(),  
    "SVM": SVC(probability=True)  
}  
for model_name, model in models.items():  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)
```

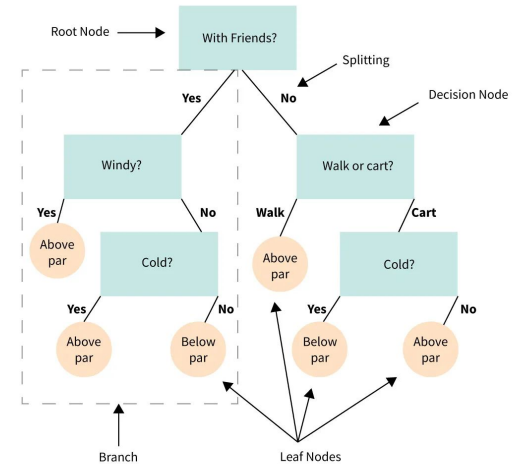
Decision Tree Classifier

Decision Trees splits data into branches based on feature values to classify data points with a sequence of actions. Structured with a root node, decision nodes, and leaf nodes.

- Tree begins with root node and splits data at each node based on feature that minimizes impurity
- At each node, the algorithm selects the split that yields the purest possible branches

Strengths: Interpretable and easy to visualize, can handle both numerical and categorical data.

Weaknesses: Prone to overfitting, and sensitive to small changes in data.



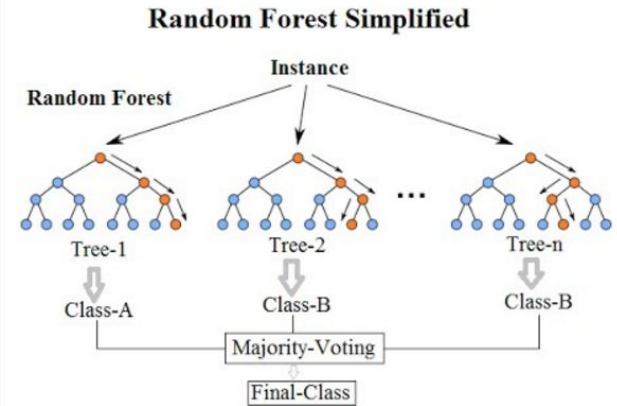
Random Forest Classifier

Random Forests combine multiple decision trees on random data subsets and combine their results for a more generalized output

- Each tree is trained on a random sample of the data
- Random feature selection occurs for each split point
- Predictions from all trees are averaged

Strengths: Reduces overfitting by averaging multiple models, and works well with high-dimensional data

Weaknesses: Computationally intensive, low interpretability




Logistic Regression

Logistic Regression is a linear model for classification, usually effective in binary cases. It estimate the probability of each class with a sigmoid function.

- Fits a linear equation to predict the probability of a data point belonging to a class
- Uses sigmoid to transform predictions into probabilities

Strengths: Simple and effective for linearly separable data.

Weaknesses: Can't model nonlinear relationships and sensitive to collinearity and outliers.




Naive Bayes

Naive Bayes is based on **Bayes' theorem**, and assumes all features are independent given the class label.

- Calculates probability of each class based on each feature independently
- Assumes independence between features
- Predictions are based on the class with the highest probability.

Strengths: Fast and efficient


Weaknesses: Independence assumption can limit accuracy by a LOT.



Support Vector Machine (SVM)

Support Vector Machine (SVM) finds a hyperplane that best separates classes by maximizing margin between classes' boundary points.


- For linearly separable data SVM finds a linear boundary
- For non-linear data, SVM uses kernels to project data into higher dimensions.




Strengths: Effective in high dimensional spaces.

Flexible with various kernel functions; allows for both linear and non-linear classification


Weaknesses: Computationally intensive with large datasets, less interpretable




Accuracy Comparison Across Datasets




ROC-AUC Comparison Across Datasets



False Positive Rate (Class 1) Comparison Across Datasets



SVM Confusion Matrix (Dataset 5)



Discussion

- **Random Forest** performed best, likely due to the ensemble approach which reduced the variance and captured more complex relationships
- **Decision Tree** was decent, but it is prone to overfitting in tasks like these due to the high dimensionality of the data
- **Logistic Regression** and **Naive Bayes** performed poorly due to underlying assumptions of linearity and feature independence, which clearly doesn't hold for this dataset.
- Future work can focus on hyperparameter tuning, further analyzing those non-linear features, and most importantly, introducing neural network models to significantly improve model performance and address class imbalance.

Thanks!